



Exercise 3.3: Finish Cluster Setup

1. View the available nodes of the cluster. It can take a minute or two for the status to change from NotReady to Ready. The NAME field can be used to look at the details. Your node name will be different.

```
student@master:~$ kubectl get node
```

```
1 NAME      STATUS  ROLES    AGE   VERSION
2 worker    Ready   <none>    50s   v1.19.1
3 master    Ready   master    28m   v1.19.1
```

2. Look at the details of the node. Work line by line to view the resources and their current status. Notice the status of Taints. The master won't allow non-infrastructure pods by default for security and resource contention reasons. Take a moment to read each line of output, some appear to be an error until you notice the status shows False.

```
student@master:~$ kubectl describe node master
```

```
1 Name:          master
2 Roles:         master
3 Labels:        beta.kubernetes.io/arch=amd64
4                beta.kubernetes.io/os=linux
5                kubernetes.io/hostname=master
6                node-role.kubernetes.io/master=
7 Annotations:   kubeadm.alpha.kubernetes.io/cri-socket: /var/run/dockerhim.sock
8                node.alpha.kubernetes.io/ttl: 0
9                projectcalico.org/IPV4Address: 10.142.0.3/32
10               volumes.kubernetes.io/controller-managed-attach-detach: true
11 CreationTimestamp: Mon, 07 Jan 2020 22:04:03 +0000
12 Taints:        node-role.kubernetes.io/master:NoSchedule
13 <output_omitted>
```

3. Allow the master server to run non-infrastructure pods. The master node begins tainted for security and performance reasons. We will allow usage of the node in the training environment, but this step may be skipped in a production environment. Note the **minus sign (-)** at the end, which is the syntax to remove a taint. As the second node does not have the taint you will get a not found error.

```
student@master:~$ kubectl describe node | grep -i taint
```

```
1 Taints:        node-role.kubernetes.io/master:NoSchedule
2 Taints:        <none>
```

```
student@master:~$ kubectl taint nodes \
--all node-role.kubernetes.io/master-
```

```
1
2 node/master untainted
3 error: taint "node-role.kubernetes.io/master:" not found
```

4. Now that the master node is able to execute any pod we **may** find there is a new taint. This behavior began with v1.12.0, requiring a newly added node to be enabled. It has reappeared in versions since then. View, then remove the taint if present. It can take a minute or two for the scheduler to deploy the remaining pods.

```
student@master:~$ kubectl describe node | grep -i taint
```

```
1 Taints:        node.kubernetes.io/not-ready:NoSchedule
2 Taints:        <none>
```

```
student@master:~$ kubectl taint nodes \
--all node.kubernetes.io/not-ready-
```

```
1
2 node/lfs58-node-1a0a untainted
3 error: taint "node.kubernetes.io/not-ready:" not found
```

5. Determine if the DNS and Calico pods are ready for use. They should all show a status of Running. It may take a minute or two to transition from Pending.

```
student@master:~$ kubectl get pods --all-namespaces
```

```
1 NAMESPACE      NAME                                     READY   STATUS    RESTARTS   AGE
2 kube-system    calico-node-jlgwr                      1/1     Running   0           6m
3 kube-system    calico-kube-controllers-74b888b647-wlqf5 1/1     Running   0           6m
4 kube-system    calico-node-tpvnr                      2/2     Running   0           6m
5 kube-system    coredns-78fcd6894-nc5cn                1/1     Running   0           17m
6 kube-system    coredns-78fcd6894-xs96m                1/1     Running   0           17m
7 <output_omitted>
```

6. **Only if** you notice the coredns- pods are stuck in ContainerCreating status you may have to delete them, causing new ones to be generated. Delete both pods and check to see they show a Running state. Your pod names will be different.

```
student@master:~$ kubectl get pods --all-namespaces
```

```
1 NAMESPACE      NAME                                     READY   STATUS             RESTARTS   AGE
2 kube-system    calico-node-qkvzh                      2/2     Running            0           59m
3 kube-system    calico-node-vndn7                      2/2     Running            0           12m
4 kube-system    coredns-576cbf47c7-rn6v4              0/1     ContainerCreating  0            3s
5 kube-system    coredns-576cbf47c7-vq5dz              0/1     ContainerCreating  0           94m
6 <output_omitted>
```

```
student@master:~$ kubectl -n kube-system delete \
pod coredns-576cbf47c7-vq5dz coredns-576cbf47c7-rn6v4
```

```
1 pod "coredns-576cbf47c7-vq5dz" deleted
2 pod "coredns-576cbf47c7-rn6v4" deleted
```

7. When it finished you should see a new tunnel, tunl0, interface. It may take up to a minute to be created. As you create objects more interfaces will be created, such as cali interfaces when you deploy pods, as shown in the output below.

```
student@master:~$ ip a
```

```
1 <output_omitted>
2 4: tunl0@NONE: <NOARP,UP,LOWER_UP> mtu 1440 qdisc noqueue state
3 UNKNOWN group default qlen 1000
4   link/ipip 0.0.0.0 brd 0.0.0.0
5   inet 192.168.0.1/32 brd 192.168.0.1 scope global tunl0
6       valid_lft forever preferred_lft forever
7 6: calib0b93ed4661@if4: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu
8 1440 qdisc noqueue state UP group default
9   link/ether ee:ee:ee:ee:ee:ee brd ff:ff:ff:ff:ff:ff link-netnsid 1
10   inet6 fe80::ecee:eeff:feee:eeee/64 scope link
11       valid_lft forever preferred_lft forever
12 <output_omitted>
```